Published using Google Docs Report abuse Learn more

JD_Sloverminds_AI_Engineering_Intern.docx

Updated automatically every 5 minutes

AI Engineering Intern

Location: Chennai, India **Duration:** 6 months (extendable) **Stipend**: 30,000

Positions: 2 interns **Preferred Institutions:** IITs, NITs, BITS, or equivalent

About Project

Solverminds builds mission-critical maritime ERP software serving global container shipping lines. Our platform manages the entire container lifecycle — booking, documentation, finance, scheduling, equipment, compliance — across 1,000+ modules and 17 business verticals.

We are building an AI Hive Mind Agency — a system of 1,000+ specialised AI agents that automate container shipping operations in real-time. The platform combines fine-tuned open-source LLMs, discrete rule engines, Neo4j knowledge graphs, and event-driven architecture. This is production software serving real shipping lines, not a research prototype.

The Role

You will build AI agents for our maritime ERP and fine-tune open-source LLMs to power them. The agents span 16 business domains — from booking validation to schedule recovery to dangerous goods compliance. Each agent is a software component with a specific job, guardrails to keep it safe, and integration with the live ERP system.

This is hands-on engineering work. You will train models, build agents, wire them into a production system, and see them run on real shipping operations.

What You'll Work On

Published using Google Docs Report abuse Learn more

- Prepare domain-specific training datasets from shipping manuals, operational Q&A pairs, and ERP transaction logs
- Run QLoRA / LoRA fine-tuning pipelines on NVIDIA RTX A6000 (49GB VRAM) — manage adapters, evaluate quality, export quantised GGUF models for production deployment
- Build evaluation frameworks — measure accuracy, hallucination rate, domain adherence, and action execution success across different model versions
- Implement continuous learning pipelines — feedback ingestion, incremental training, and model versioning without full retraining

Building AI Agents with Guardrails

- Build micro-agents (discrete, hybrid, and LLM-based) that handle field-level validations, transaction execution, and cross-domain monitoring across the ERP
- Implement the 3-level agent hierarchy: Domain Directors (orchestrators) → Screen Managers (coordinators) → Micro-Agents (workers) — each with clear scope and guardrails
- Build guardrails for every agent: input validation, output verification, confidence thresholds, fallback to human review, and audit logging of every agent action
- Implement the event-driven publish/subscribe architecture — user actions fire events, agents respond in parallel, screen manager collects and merges results within latency bounds
- Build doer agents that execute ERP transactions (booking creation, BL issuance, invoice posting) with pre-execution validation and rollback capability
- Build follow-up agents that monitor cross-domain completion — detect process gaps (BL issued but no invoice, vessel sailed but no CAR) and trigger corrective actions
- Design agent communication protocols — how agents publish results, subscribe to events, handle conflicts, and escalate to human operators when confidence is low

Knowledge Graph Integration

Published using Google Docs Report abuse Learn more

JD_Sloverminds_AI_Engineering_Intern.docx Updated automatically every 5 minutes

system smarter over time
- Write Cypher queries for agent consumption — path traversal, pattern matching, cascade impact analysis

What We're Looking For

Must Have

- Strong Python programming — you'll write a lot of it, and it needs to be production quality
- Understanding of how LLMs work — attention, tokenisation, inference, context windows, temperature, sampling
- Hands-on experience with at least one LLM framework: Hugging Face Transformers, Ollama, vLLM, LangChain, or similar
- Understanding of REST APIs (FastAPI preferred) — agents expose and consume APIs
- Ability to read and understand existing codebases — you'll work within a large maritime ERP
- Experience fine-tuning LLMs (LoRA, QLoRA, PEFT, Unsloth)
- Experience with agent frameworks (CrewAI, AutoGen, LangGraph, or similar)
- Knowledge of Neo4j or any graph database
- Experience with event-driven architectures (Kafka, Redis pub/sub)
- Familiarity with embeddings, vector search, and RAG patterns
- Experience with NVIDIA GPU workflows (CUDA, mixed precision training)
- Spring Boot / Java exposure (the ERP backend is Java-based)

What Matters Most

- You build things that work, not just prototype — your code will go into production
- You're comfortable with ambiguity — this is new territory, not a textbook assignment
- You understand that "agent" means a software component with a specific job, guardrails, and accountability — not a chatbot

Published using Google Docs Report abuse Learn more

JD_Sloverminds_AI_Engineering_Intern.docx Updated automatically every 5 minutes

that escalates to a human

What You'll Learn

- End-to-end LLM fine-tuning for production — from dataset preparation through training, evaluation, quantisation, and deployment on enterprise GPU hardware
- How to build production multi-agent systems at scale — 1,000+ agents, event-driven, with real-time latency requirements
- Agent guardrail design — input validation, output verification, confidence scoring, human escalation, audit trails
- Neo4j knowledge graphs for AI reasoning — how agents use graph context to make better decisions
- Maritime shipping domain knowledge — booking, BL, D&D, stowage, scheduling — transferable to any complex logistics domain
- The discipline of shipping AI to production — versioning, monitoring, evaluation, and continuous improvement

Work Environment

- You'll work alongside experienced engineers on production systems, not isolated intern projects
- The maritime shipping domain is complex — expect a learning curve, and embrace it
- Core technologies: Python, FastAPI, Gemma/Mistral/Llama (fine-tuned), Neo4j, Kafka, Claude (Anthropic)
- Hardware: NVIDIA RTX A6000 (49GB VRAM) for training and inference
- Code quality matters — your code goes into production systems used by global shipping lines
- We value people who build things that work over people who know theory but can't ship

What We Offer

- Real production AI engineering — not a made-up intern project

Published using Google Docs Report abuse Learn more

JD_Sloverminds_AI_Engineering_Intern.docx Updated automatically every 5 minutes

performers

Please Note :

1. **The last date to apply for this opportunity is April 10, 2026 by 11 AM IST.**
2. **Not adhering to placement and internship policy rules will lead to necessary disciplinary actions by the committee.**